

AJAY KAVITI

kaviay0408@gmail.com | (216) 566-4055 | [linkedin.com/in/ajay-kaviti/](https://www.linkedin.com/in/ajay-kaviti/) | github.com/ajaykaviti | Naperville, IL

PROFESSIONAL SUMMARY

Senior Cloud Data Engineer with 5+ years designing and operating cloud-native data platforms on Snowflake, AWS, and Azure following medallion architecture and data lakehouse principles. Proven record cutting pipeline runtimes by up to 60%, eliminating vendor licensing costs, and migrating legacy ETL and MDM workloads to modern ELT architectures using dbt, Airbyte, Fivetran, Prefect, Airflow, PySpark, and Databricks. Experienced in data lineage tracking, cost optimization, and enforcing data governance across multi-cloud environments. Certified SnowPro Core and Microsoft Azure DP-203. M.S. Information Technology, University of Cincinnati.

TECHNICAL SKILLS

Data Platforms:	Snowflake, Databricks, Redshift, Azure Synapse, DynamoDB, PostgreSQL, MySQL, Delta Lake, Apache Iceberg
Ingestion & ELT:	Airbyte, Fivetran, dbt, Snowflake Streams & Tasks, Dynamic Tables, data lakehouse, medallion architecture
Orchestration:	Prefect, Apache Airflow, AWS Step Functions, Azure Data Factory, AWS Glue, data lineage, data mesh
Cloud & Infra:	AWS (S3, Glue, Kinesis, ECS Fargate, Lambda, QuickSight), Azure (DevOps, ADF, Synapse), cost optimization
Languages:	Python, PySpark, SQL, Java, T-SQL
DevOps & IaC:	Terraform, Docker, Kubernetes, GitHub Actions, Azure DevOps, AWS CodeBuild, ECR
Visualization:	Power BI, Looker, Tableau

PROFESSIONAL EXPERIENCE

Senior Cloud Data Engineer - Oncourse Home Solutions - Naperville, IL

Dec 2024 – Present

- Architected a centralized Airbyte and Snowflake ingestion hub following medallion architecture principles (bronze, silver, gold layers), implementing plug-and-play connector templates that eliminated all custom loader code and replaced a proprietary vendor, cutting recurring licensing costs entirely.
- Converted claims data processing from nightly full-refresh loads to incremental ELT using Airbyte syncs and Snowflake Streams and Tasks, slashing pipeline runtime from 5 hours to 2 hours (60% reduction) with proportional compute savings.
- Led zero-data-loss migration of a legacy MDM SQL Server warehouse to Snowflake, refactoring 80+ stored procedures with SnowSQL and dbt and coordinating cut-over scheduling through Azure DevOps release pipelines.
- Engineered layered dbt models across staging, intermediate, and mart layers with column-level tests, snapshots, source freshness checks, and data lineage tracking, improving overall data reliability and reducing analyst-reported incidents.
- Optimized Snowflake warehouse costs through dynamic scaling, multi-cluster queues, auto-clustering on high-cardinality columns, and nightly auto-suspend, achieving measurable cost optimization without impacting SLAs.
- Delivered end-to-end Azure DevOps CI/CD pipelines that lint, unit-test, and deploy dbt, Prefect, and Terraform code with automated blue-green releases and one-click rollback, cutting deployment risk significantly.
- Automated secrets rotation and Snowflake RBAC policy enforcement across Azure DevOps environments, producing audit-ready access evidence on demand and reducing manual compliance effort.
- Applied Amazon S3 intelligent-tiering lifecycle policies and query result caching across Snowflake workloads, optimizing storage costs and accelerating repeated dashboard queries by up to 40%.

Data Engineer (Contract) - CenturyLink - Remote

Aug 2024 – Nov 2024

- Designed and centralized Fivetran-to-Snowflake ingestion pipelines, standardizing all connector configurations and eliminating ad-hoc custom loader scripts maintained across the data team.
- Refactored legacy SQL-based transformation logic into layered dbt models across staging, intermediate, and mart layers with data tests and snapshots, reducing maintenance overhead by consolidating transforms into version-controlled code.
- Converted batch-heavy pipelines to incremental Snowflake Streams and Tasks models, reducing end-to-end processing time by 35% and significantly lowering warehouse compute consumption.
- Implemented Snowflake Dynamic Tables to automate incremental transformations and manage dependency chains, reducing manual Airflow DAG complexity and improving data freshness for analytics workloads.
- Orchestrated all data workflows using Apache Airflow, achieving 99.9%+ SLA adherence through proactive webhook-based alerting, on-call runbooks, and automated retry policies.
- Partnered with analytics and business intelligence teams to design and deliver curated Snowflake data marts powering executive-level Looker dashboards, replacing fragmented ad-hoc queries.
- Implemented Azure DevOps pipelines to lint, test, and deploy dbt and Terraform code with automated rollback, enforcing code quality standards and reducing production incidents.

AWS & Snowflake Data Engineer • Infosys • Hyderabad, India

Nov 2021 – Dec 2022

- Architected a Snowflake data lakehouse on AWS as the central analytics platform for a 10 TB retail dataset, designing medallion-layer dbt models (staging, intermediate, marts) with column-level tests and source freshness checks that reduced reporting query times by 40%.
- Orchestrated end-to-end ELT workflows using Apache Airflow, scheduling and monitoring 50+ dbt model DAGs with SLA-based alerting, automated retries, and Slack webhook notifications to maintain 99.8%+ pipeline uptime.
- Built layered dbt transformations on Snowflake incorporating incremental models, snapshots, and macros to process 500+ GB of daily retail data, cutting full-refresh runtimes by 35% and reducing warehouse compute costs.
- Engineered AWS Glue and Step Functions data ingestion pipelines extracting from 15+ source systems into Snowflake S3-based staging zones, enforcing schema validation and data quality checks at every layer.
- Implemented Snowflake Streams and Tasks for CDC-based incremental loads on high-volume transactional tables, eliminating nightly full-refresh batch jobs and improving data freshness from 24 hours to under 2 hours.
- Designed and maintained Airflow DAGs for orchestrating cross-system data dependencies between Snowflake, AWS Glue, S3, and Redshift, achieving full data lineage visibility and reducing pipeline failures by 45%.
- Developed self-serve Snowflake data marts consumed by Power BI and Amazon QuickSight dashboards, delivering analytics over 1 TB of retail data and increasing data-driven decision-making efficiency by 25%.
- Hardened Snowflake and AWS environments with role-based access control, KMS encryption, network policies, VPC endpoints, and GuardDuty threat detection, securing sensitive retail PII across all pipeline stages.

Software Associate • Dhatsol • Hyderabad, India

Aug 2019 – Oct 2021

- Engineered Amazon Kinesis Data Streams pipelines to ingest approximately 500,000 trading ticks per minute with zero downtime during live market hours, ensuring continuous data availability for downstream analytics.
- Developed Kinesis Data Analytics jobs that cleaned, enriched, and normalized real-time feeds from exchanges, news wires, and market-data vendors, significantly reducing downstream data quality issues.
- Designed table schemas in Amazon Redshift and the AWS Glue Data Catalog, choosing optimal sort and distribution keys that accelerated analyst queries and simplified ongoing schema maintenance.
- Modeled high-velocity trading data in Amazon DynamoDB and Amazon Keyspaces (for Apache Cassandra), selecting partition strategies that supported rapid point-reads and range scans under market-hours load spikes.
- Implemented response caching in Amazon API Gateway for symbol-lookup endpoints, reducing backend load and cutting average API latency for high-frequency trading clients.
- Leveraged AWS Lake Formation to govern and share S3-based datasets securely with external research partners, enforcing fine-grained column-level access controls without duplicating data.
- Engineered CI/CD pipelines with AWS CodeCommit, CodeBuild, and GitHub Actions to automate build, test, and deployment for all ETL scripts, reducing manual release effort and deployment errors.
- Instrumented CloudWatch metrics and Lambda-based data-quality checks that flagged record-count mismatches and schema drift in near-real-time, improving pipeline observability and on-call response times.

EDUCATION**M.S. Computer Engineering** — University of Cincinnati, Cincinnati, OH

Apr 2024

B.Tech Computer Engineering — KL University (KLEF), Guntur, India

May 2020

CERTIFICATIONS

- **SnowPro Core** – Snowflake • 2025 – Covers Snowflake platform architecture, virtual warehouses, data loading and transformation, performance optimization, data sharing, and cloud security best practices.
- **Azure DP-203** – Microsoft • 2024 – Validates expertise in designing and implementing data storage solutions, data processing pipelines, and security and compliance controls on Microsoft Azure.
- **dbt Certified Developer** – dbt Labs • 2026 – Validates advanced proficiency in building modular dbt projects, writing tests, managing snapshots, optimizing incremental models, and implementing data transformation best practices.
- **SQL Developer Certified** – Oracle • 2022 – Demonstrates advanced proficiency in SQL development, complex query writing, performance tuning, indexing strategies, and relational database design best practices.
- **Tableau Desktop Specialist** – Salesforce • 2022 – Confirms skills in connecting to data sources, building interactive visualizations and dashboards, applying filters and calculations, and delivering analytics insights using Tableau Desktop.